


NEVER STOP BUYING LOTTERY TICKETS,
NO MATTER WHAT ANYONE TELLS YOU.

I FAILED AGAIN AND AGAIN, BUT I NEVER
GAVE UP. I TOOK EXTRA JOBS AND
POURED THE MONEY INTO TICKETS.

AND HERE I AM, PROOF THAT IF YOU
PUT IN THE TIME, IT PAYS OFF!




EVERY INSPIRATIONAL SPEECH BY SOMEONE
SUCCESSFUL SHOULD HAVE TO START WITH
A DISCLAIMER ABOUT SURVIVORSHIP BIAS.

Lecture 10

11.2.21

The problem with multiple testing



1. If you roll a die once, what is the probability of rolling a 6?

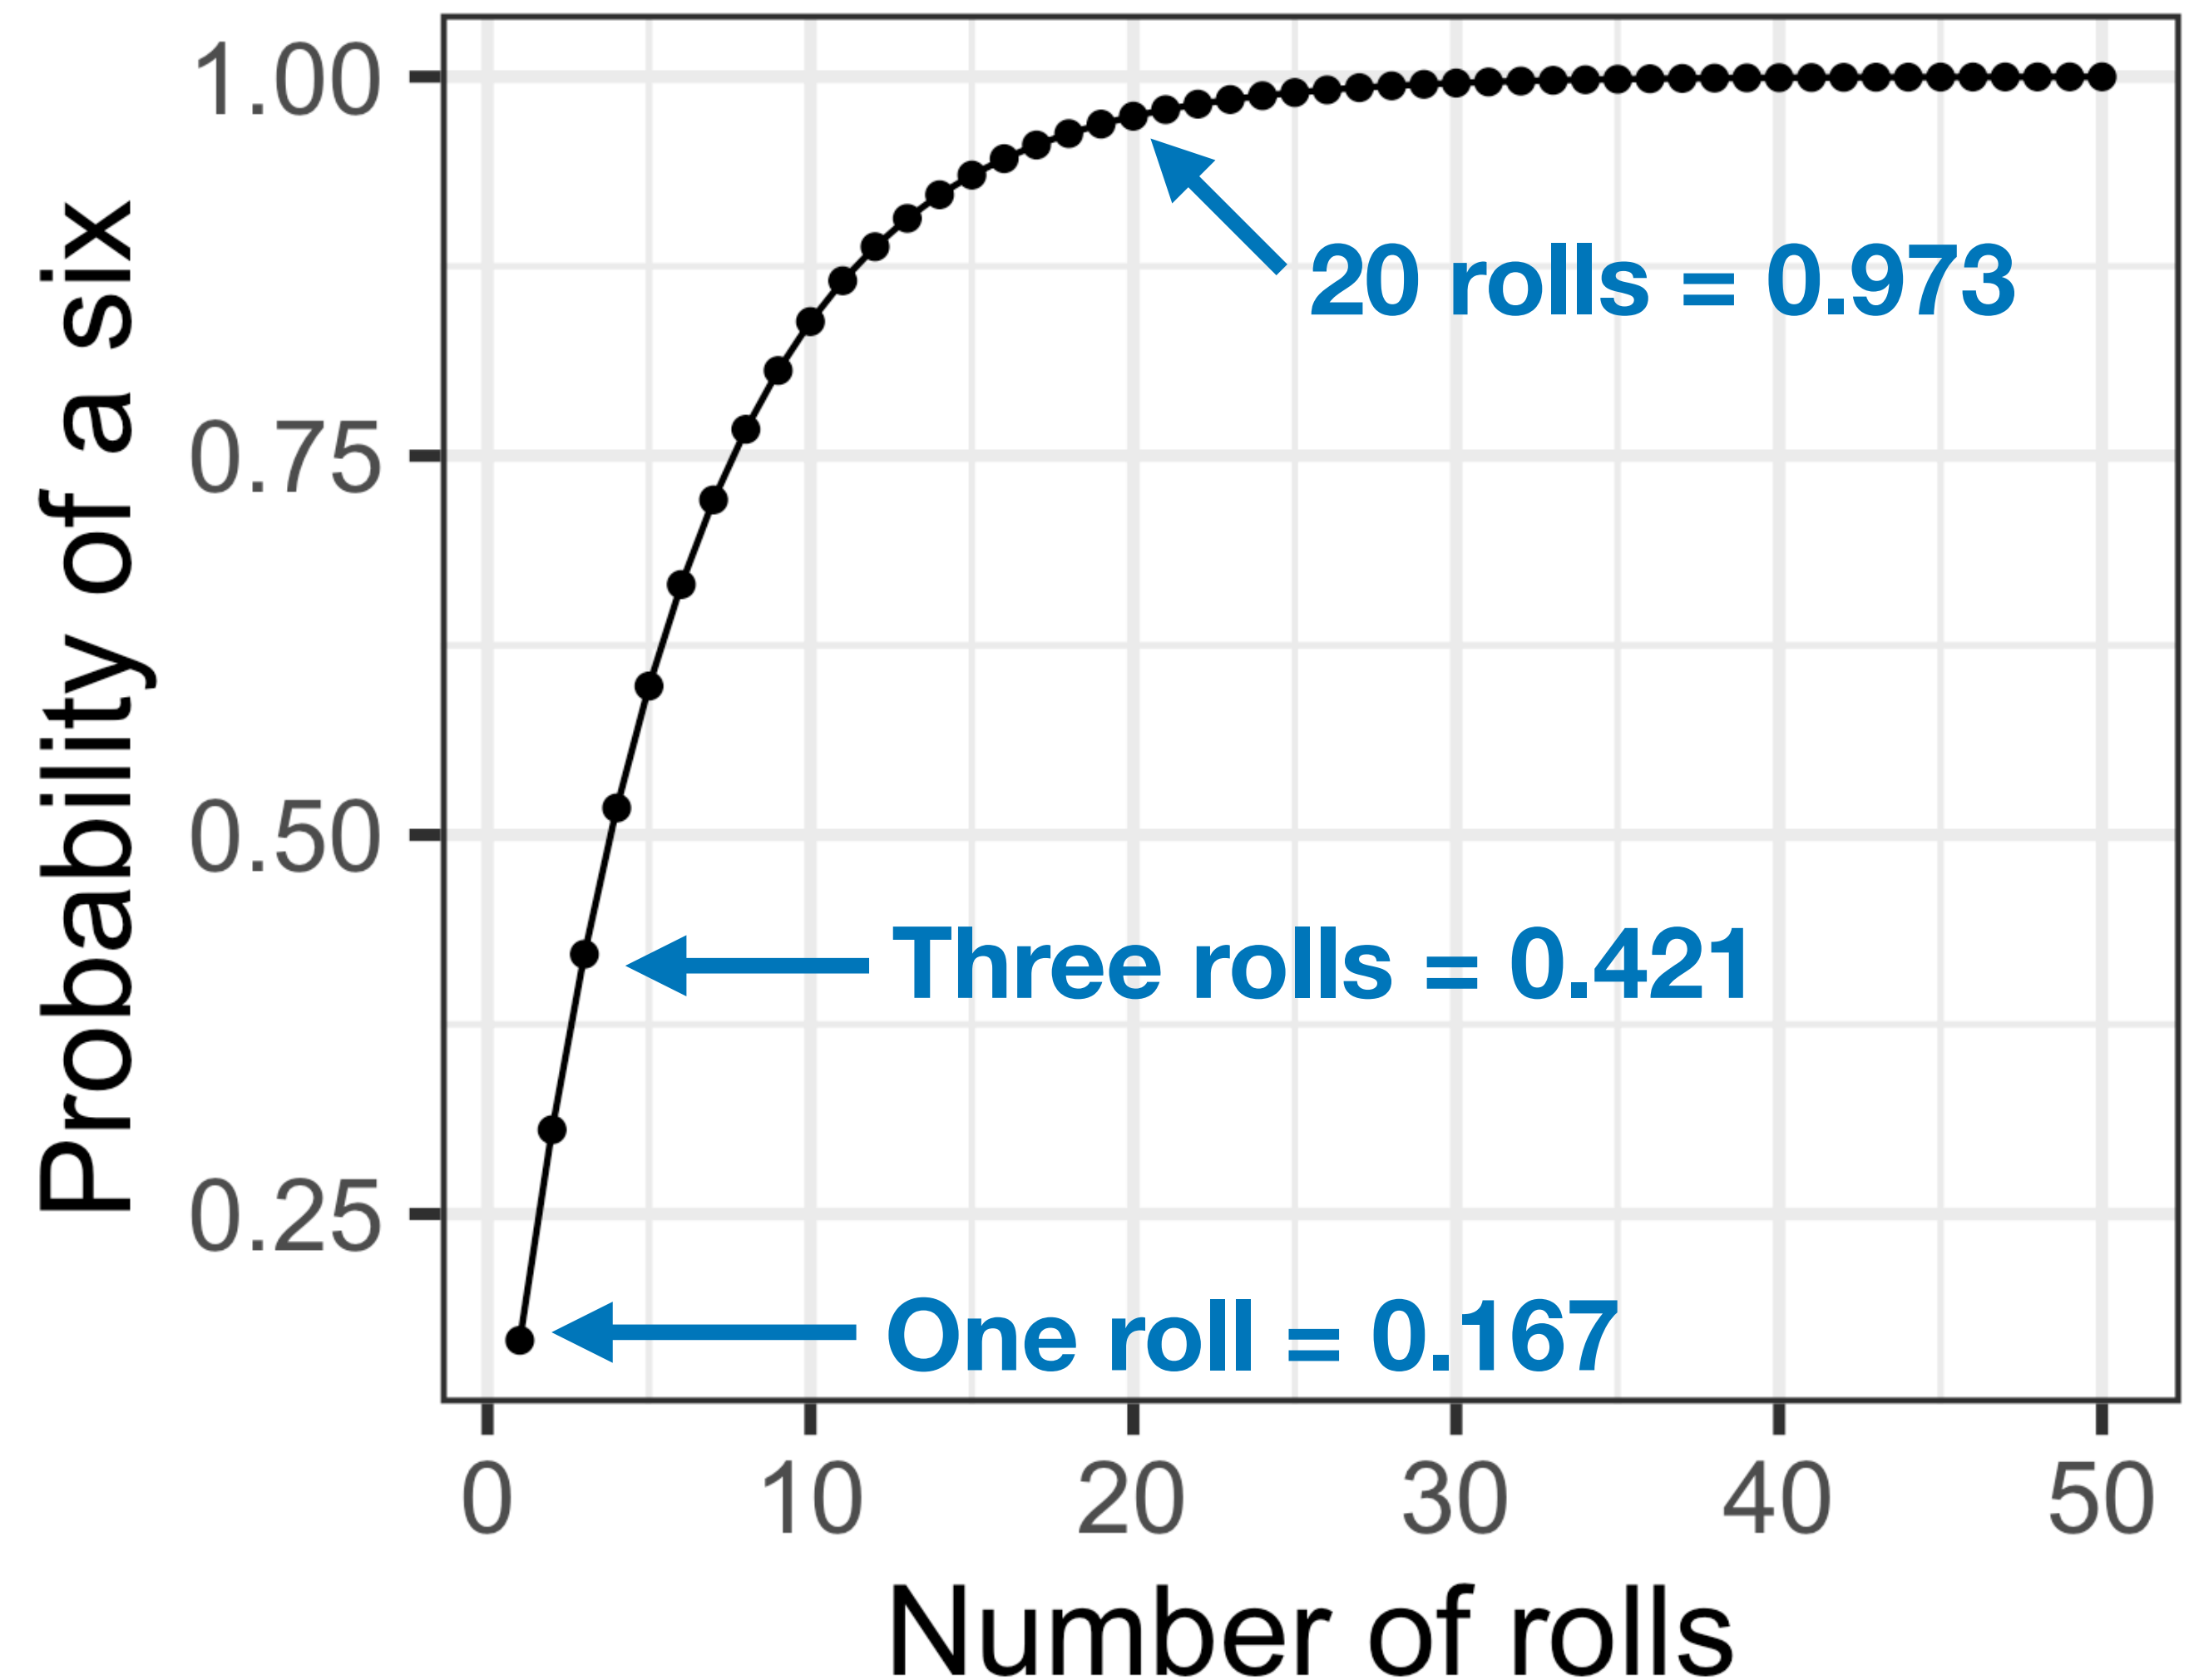

$$\text{Pr}\{\text{rolling a six}\}: 1/6 = 0.1667$$

2. If you roll a die three times, what is the probability of rolling a 6 at least once?

$$\text{Pr}\{\text{at least one six}\} = 1 - \text{Pr}\{\text{no sixes}\}$$

$$\text{Pr}\{\text{at least one six}\} = 1 - [(5/6)(5/6)(5/6)] = 0.421$$

The problem with multiple testing



The problem with multiple testing

t test, $\alpha = 0.05$:

**Remember: α is the Type I
(false positive) error rate**

One test:


$$1 - (1 - 0.05)^1 = 0.05$$

20 tests:

$$1 - (1 - 0.05)^{20} = 0.641$$

100 tests:

$$1 - (1 - 0.05)^{100} = 0.994$$



Refresher: Type I and type II errors

Null hypothesis is...	TRUE	FALSE
REJECTED	Type I error <i>(False positive)</i>	Correct! <i>(True positive)</i>
NOT REJECTED	Correct! <i>(True negative)</i>	Type II error <i>(False negative)</i>

↓ Type I Error → ↑ Type II Error

Multiple hypothesis correction

Adjusting α in some way so that the probability of a test being significant by chance is lower (remains below the significance threshold)

Bonferroni correction (FWER)

False Discovery Rate (FDR)

Family-wise error rate (FWER)

FWER = probability of at least one type I error

- Most common FWER is the **Bonferroni correction**
 - Based on the idea that the probability that at least one of several events will occur cannot exceed the sum of the individual probabilities

0.05 $\alpha = 0.05$; 3 tests; H_0 is TRUE

P(at least one test is significant) $\leq 0.05 + 0.05 + 0.05$

P(at least one test is significant) $\leq 3(0.05)$

$$0.05 \leq 3(X)$$

$$\frac{0.05}{3} \leq X$$

Bonferroni: divide α / n and use this value as the new significance threshold α

The Bonferroni correction

Knockout of 10 different genes, measured a phenotype of interest and performed a *t*-test to compare each knockout to the control. The following are the p-values obtained:

Which genes are significant at $\alpha = 0.05$?

A	B	C	D	E	F	G	H	I	J
0.084	0.036	0.063	0.186	0.108	0.042	0.01	0.132	0.175	0.0012

Which genes are significant at $\alpha = 0.05/10 = 0.005$?

The Bonferroni correction

Knockout of 10 different genes, measured a phenotype of interest and performed a *t*-test to compare each knockout to the control. The following are the p-values obtained:

Which genes are significant at $\alpha = 0.05/10 = 0.005$?

A	B	C	D	E	F	G	H	I	J
0.084	0.036	0.063	0.186	0.108	0.042	0.01	0.132	0.175	0.0012

$$1 - (1 - 0.05)^{10} = 0.401 \quad \longrightarrow \quad 1 - \left(1 - \frac{0.05}{10}\right)^{10} = 0.048$$

Bonferroni correction is very **strict/conservative**—It is based on the fact that the null hypothesis is true, and could lead to a **high percentage of false negatives**

The Bonferroni correction

- Assumes all tests are independent — although they often are not, which can lead to a high percentage of false negatives (i.e. Type II errors)
- Counter-intuitive: the interpretation of findings depends on the number of other tests run simultaneously
- Simple to perform and does reduce the false positives (i.e. Type I errors)

Bonferroni correction is very **strict/conservative**—It is based on the fact that the null hypothesis is true, and could lead to **a high percentage of false negatives**

Multiple hypothesis correction

Bonferroni correction (FWER)

- Easy to perform
- Very strict/conservative
- Set threshold to α/n
- You care more about preventing false positives than false negatives
- Sample size is high
- Effect of interest is very large and/or consistent

False Discovery Rate (FDR)

False discovery rate (FDR)

FDR = expected proportion of false positives among all significant results

Null hypothesis is...	TRUE	FALSE
REJECTED	Type I error (<i>False positive</i>)	Correct! (<i>True positive</i>)
NOT REJECTED	Correct! (<i>True negative</i>)	Type II error (<i>False negative</i>)

$$\text{FDR} = \frac{\text{Number of false positives}}{\text{False positives} + \text{true positives}}$$

False discovery rate (FDR)

FDR = expected proportion of false positives among all significant results

Null hypothesis is...	TRUE	FALSE
REJECTED	Type I error (False positive)	Correct! (True positive)
NOT REJECTED	Correct! (True negative)	Type II error (False negative)

$$\text{FDR} = \frac{\text{Number of false positives}}{\text{False positives} + \text{true positives}}$$

$$\text{FPR}^* = \frac{\text{Number of false positives}}{\text{False positives} + \text{true negatives}}$$

**false positive rate*

False discovery rate (FDR)

FDR = expected proportion of false positives among all significant results

- Benjamini-Hochberg method is a popular form of FDR
 - If many variables are significant, then surely there must be a true effect
 - Therefore, the actual chance of a false positive is much lower than FWER suggests
- FDR estimates the rejection region so that $FDR < \alpha$

$$FDR = \frac{\text{Number of false positives}}{\text{False positives} + \text{true positives}}$$

*J. R. Statist. Soc. B (1995)
57, No. 1, pp. 289-300*

**Controlling the False Discovery Rate: a Practical and Powerful
Approach to Multiple Testing**


By YOAV BENJAMINI† and YOSEF HOCHBERG

Tel Aviv University, Israel


An aside on p -value distributions

Suppose we are testing many gene's expressions before and after drug treatment

Most of the time, genes will not be expressed differently




Sometimes, genes **will** be expressed differently




Remember: a p -value of 0.05 means 5% of the time you will get this value by chance

An aside on p -value distributions

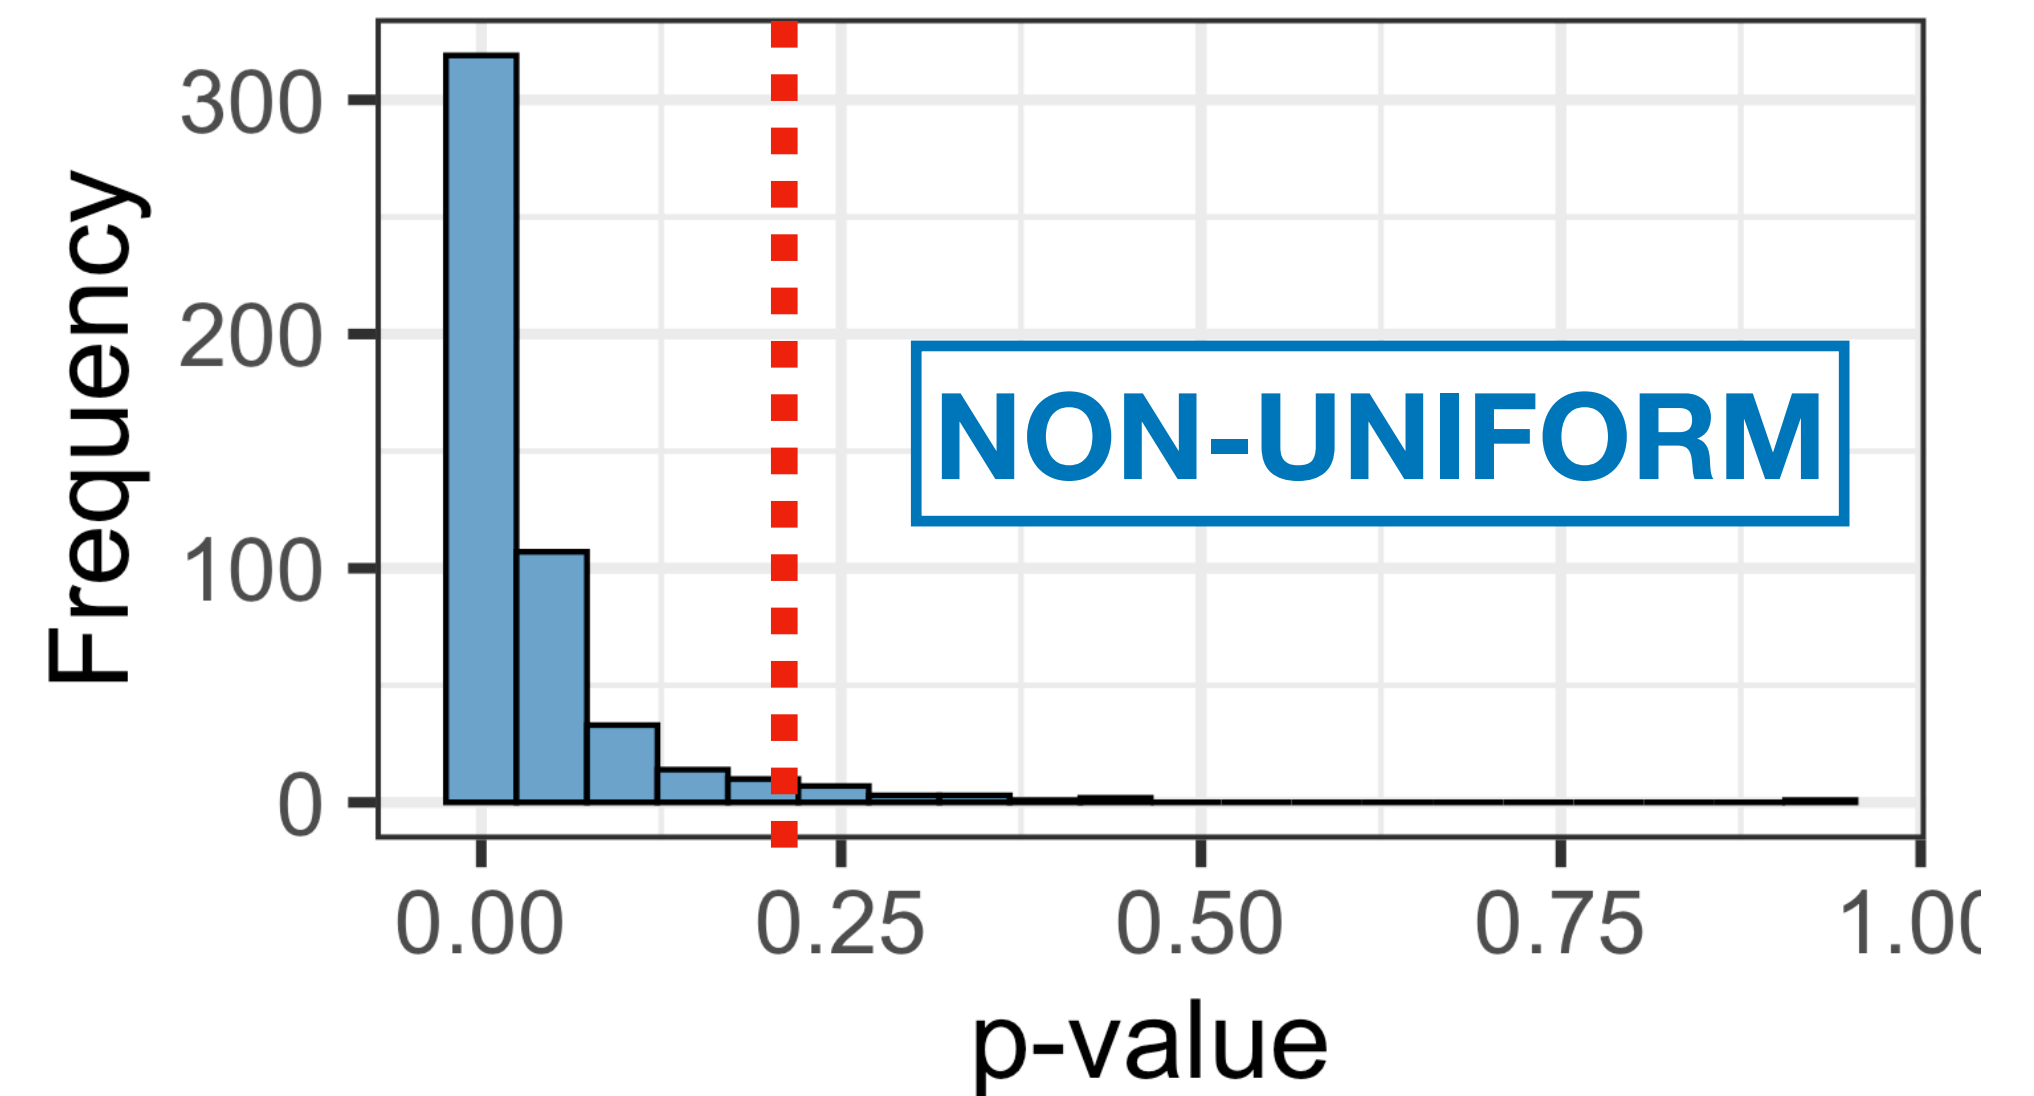

Suppose we are testing many gene's expressions before and after drug treatment



Most of the time, genes will not be expressed differently




Sometimes, genes **will** be expressed differently




Remember: a p -value of 0.05 means 5% of the time you will get this value by chance

An aside on p -value distributions


Suppose we are testing many gene's expressions before and after drug treatment



Most of the time, genes will not be expressed differently




Sometimes, genes **will** be expressed differently



Adjusts p-values (makes them larger) so that **5% of the “significant” results** will be **false positives**.

An aside on p -value distributions



**not to scale*

Calculating the FDR

FDR = expected proportion of false positives among all significant results

A	B	C	D	E	F	G	H	I	J
0.084	0.036	0.063	0.186	0.108	0.042	0.024	0.132	0.175	0.0012

1. Order p-values from smallest to largest

Calculating the FDR

FDR = expected proportion of false positives among all significant results

J	G	B	F	C	A	E	H	I	D
0.0012	0.024	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186

1. Order p-values from smallest to largest



2. Rank the p-values

Calculating the FDR

FDR = expected proportion of false positives among all significant results

J	G	B	F	C	A	E	H	I	D
0.0012	0.024	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186
1	2	3	4	5	6	7	8	9	10

1. Order p-values from smallest to largest



2. Rank the p-values



3. Adjust the p-values (starting with largest):

Current p-value * (total number of p-values)/(rank of p-value)

Calculating the FDR

FDR = expected proportion of false positives among all significant results

J	G	B	F	C	A	E	H	I	D
0.0012	0.024	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186
1	2	3	4	5	6	7	8	9	10
									0.186

1. Order p-values from smallest to largest ✓

2. Rank the p-values ✓

*Largest p-value is
always the same*

3. Adjust the p-values (starting with largest):

Current p-value * (total number of p-values)/(rank of p-value)

Calculating the FDR

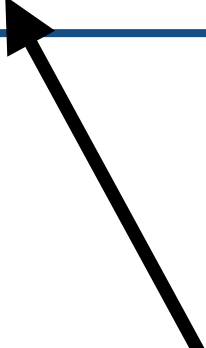
FDR = expected proportion of false positives among all significant results

J	G	B	F	C	A	E	H	I	D
0.0012	0.024	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186
1	2	3	4	5	6	7	8	9	10
								0.186	0.186

1. Order p-values from smallest to largest 

2. Rank the p-values 

$0.175(10/9) = 0.194$



3. Adjust the p-values (starting with largest):

Previous p-value **OR** Current p-value * (total number of p-values)/(rank of p-value)

Calculating the FDR

FDR = expected proportion of false positives among all significant results

J	G	B	F	C	A	E	H	I	D
0.0012	0.01	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186
1	2	3	4	5	6	7	8	9	10
0.012	0.05	0.12	0.105	0.126	0.14	0.154	0.165	0.186	0.186

$$\alpha = 0.05$$

1. Order p-values from smallest to largest ✓

2. Rank the p-values ✓

3. Adjust the p-values (starting with largest): ✓

Previous p-value **OR** Current p-value * (total number of p-values)/(rank of p-value)

Multiple hypothesis correction

Bonferroni correction (FWER)

- Easy to perform
- Very strict/conservative
- Set threshold to α/n
- You care more about preventing false positives than false negatives
- Sample size is high
- Effect of interest is very large and/or consistent

False Discovery Rate (FDR)

- More complicated to do
- More permissive
- Set threshold so 5% of positives are false
- Both minimize false positives and keep false negatives low
- Limited sample size
- Effect of interest is not very large nor consistent

Comparing Bonferroni vs. FDR

	J	G	B	F	C	A	E	H	I	D
p-value	0.0012	0.01	0.036	0.042	0.063	0.084	0.108	0.132	0.175	0.186
$\alpha < 0.05$	TRUE	TRUE	TRUE	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
BF	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
FDR	TRUE	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE

- Very significant cases (i.e. gene “J”) will always be significant
- Border cases are often the most interesting (and common) in biology
- One may use **FDR** to explore and **Bonferroni/FWER** to confirm

Bonferroni and FDR correction in R

```
# create vector of p-values
> vals <- c(0.0012, 0.01, 0.036, 0.042, 0.063,
0.084, 0.108, 0.132, 0.175, 0.186)

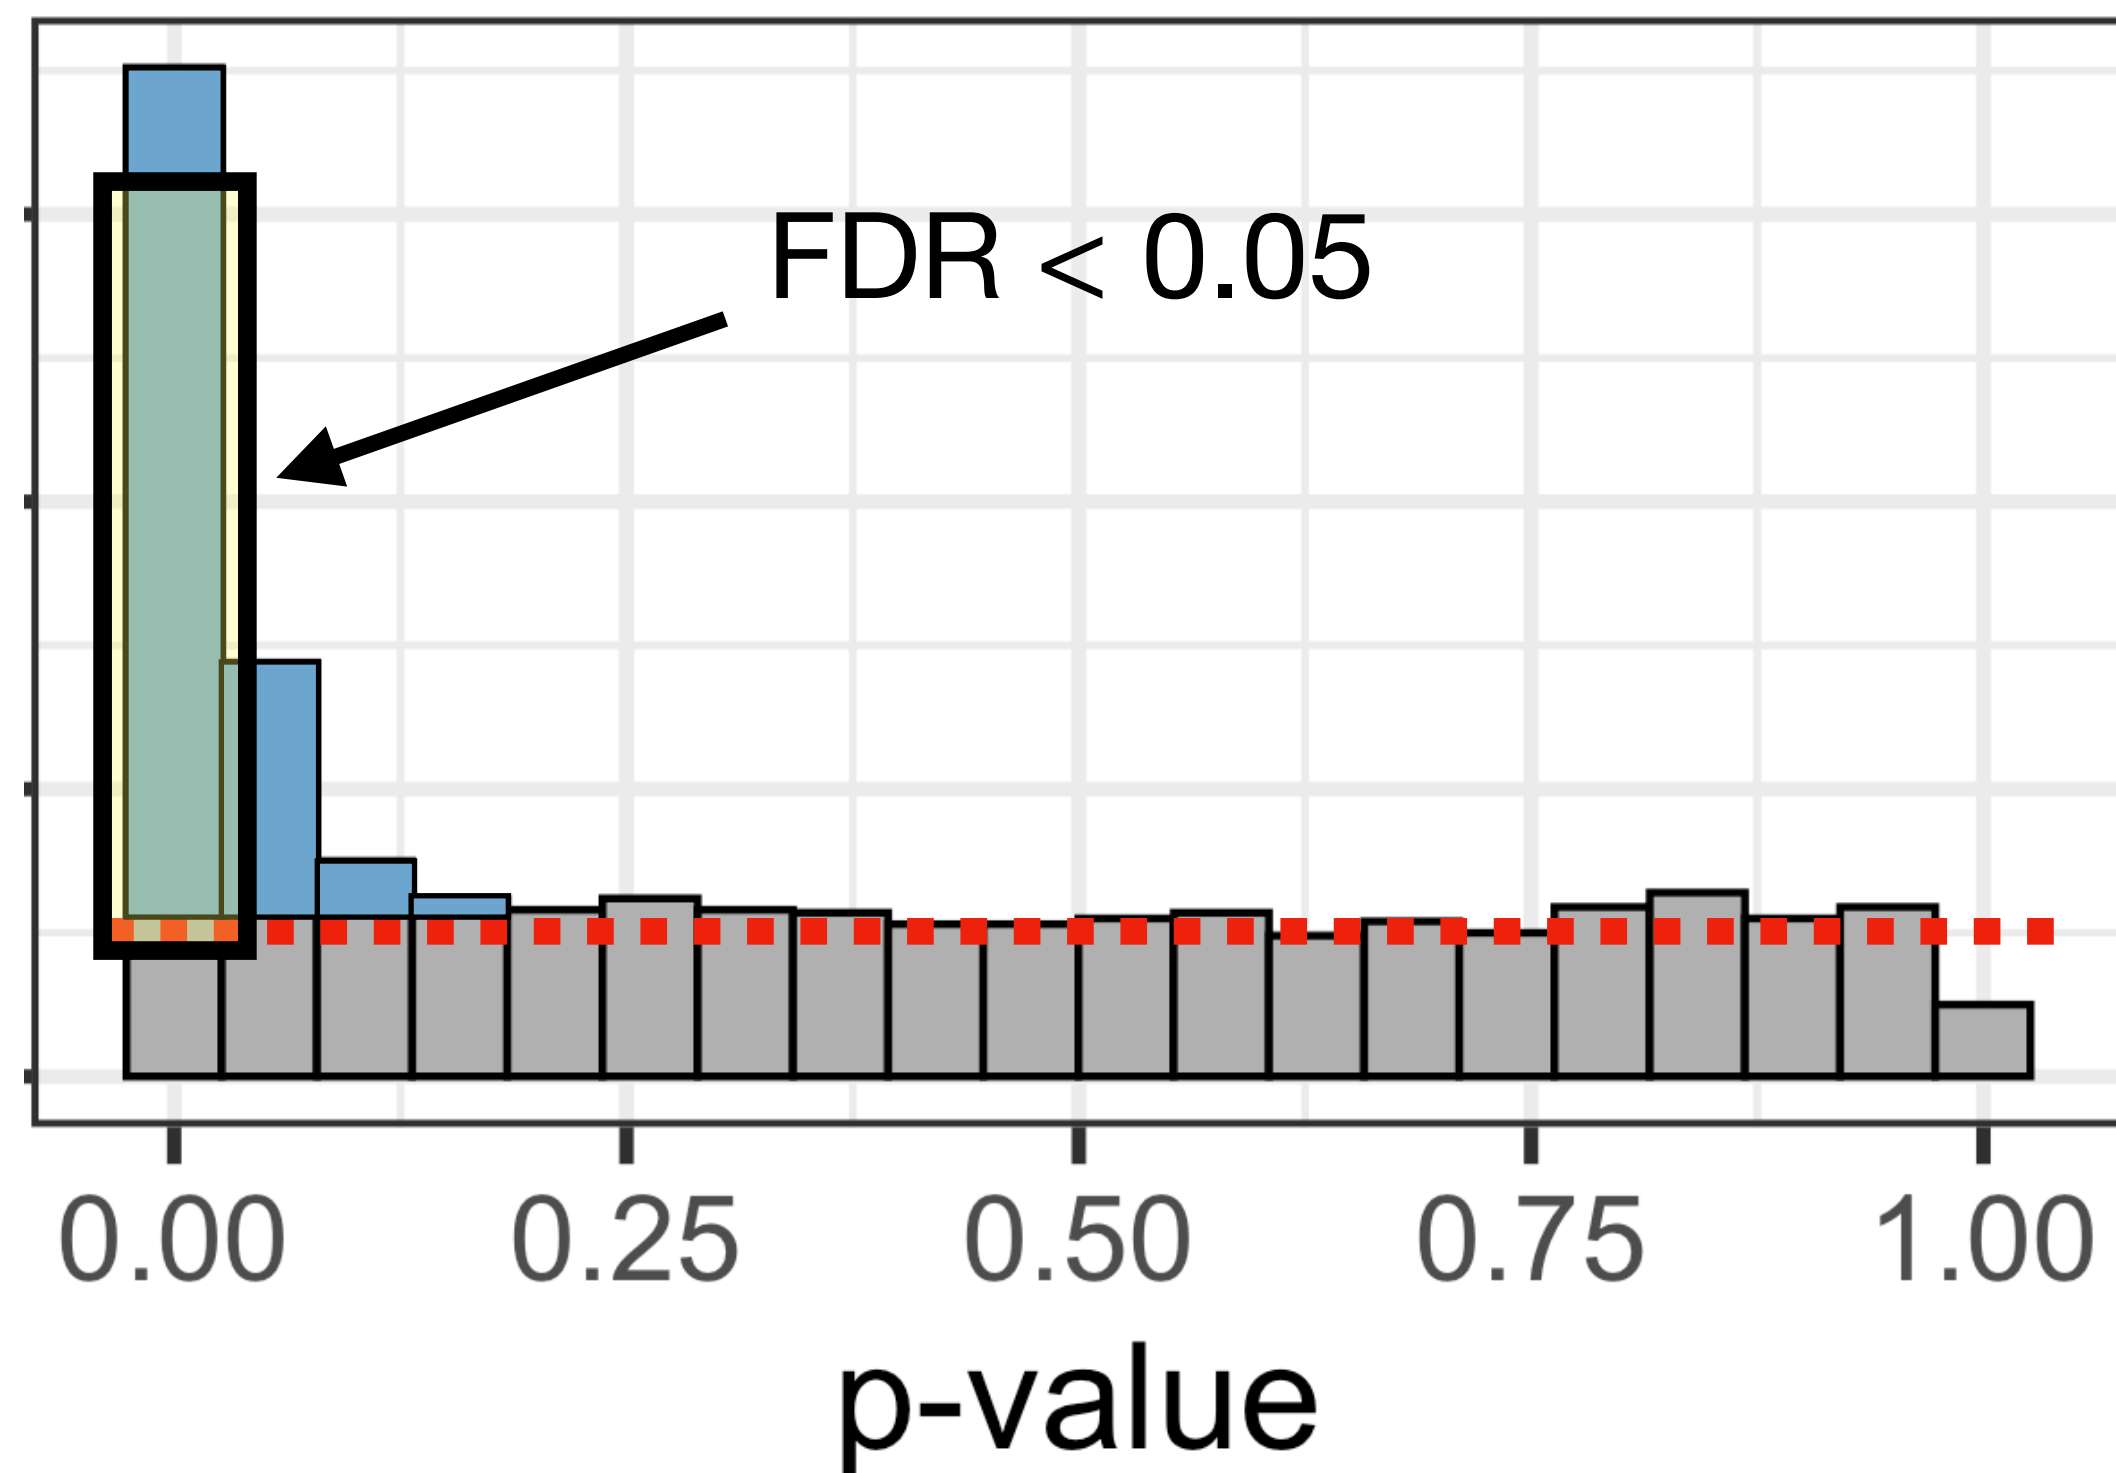
# adjust p-values with bonferroni correction
> p.adjust(vals, method = "bonferroni")

[1] 0.012 0.100 0.360 0.420 0.630 0.840 1.000 1.000
1.000 1.000

# adjust p-values with FDR correction
> p.adjust(vals, method = "fdr")

[1] 0.01200 0.05000 0.10500 0.10500 0.12600 0.14000
0.15428 0.16500 0.18600 0.18600
```

Bonferroni and FDR correction in R



$$\text{FDR} = \frac{\text{Number of false positives}}{\text{False positives} + \text{true positives}}$$

Table of unadjusted p-values
that are significant

	FALSE	TRUE
Two-distributions	114	386
Uniform	957	43

```
> fdr <- 43 / (43 + 386)
```

Table of FDR-adjusted p-values
that are significant

	FALSE	TRUE
Two-distributions	350	150
Uniform	995	5

```
> fdr <- 5 / (5 + 150)
```

Beyond Bonferroni and FDR

- Adjusting the p-value to reduce number of false positives is a very active area of statistics
- There are many more ways, but these two are perhaps the most common, especially in biology
- Others: positive false discovery rate (pFDR), Holm (type of FWER), local false discovery rate (local FDR), Fisher's least significant difference (Fisher's LSD), Tukey's honest significant difference (Tukey's HSD), permutation/randomization

Beyond Bonferroni and FDR

- Adjusting the p-value to reduce number of false positives is a very active area of statistics
- There are many more ways, but these two are perhaps the most common, especially in biology
- Others: positive false discovery rate (pFDR), Holm (type of FWER), local false discovery rate (local FDR), Fisher's least significant difference (Fisher's LSD), Tukey's honest significant difference (Tukey's HSD), **permutation/randomization**

Permutation and FWER


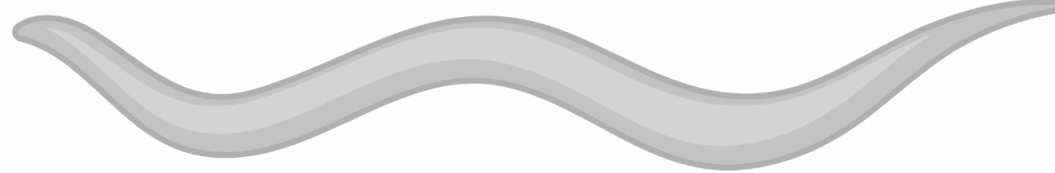
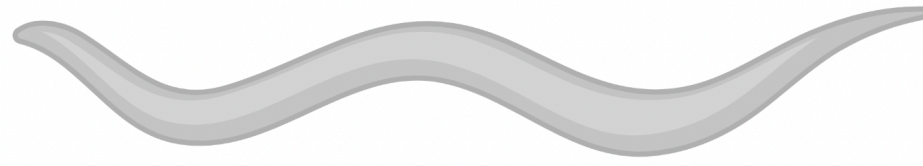



- One main disadvantage for Bonferroni FWER correction is that it assumes all tests are independent, when often they are correlated
- Permutation, by definition, tests the null hypothesis that there is no effect and breaks down any correlation structure of the data
- **(1) Randomize the data and conduct hypothesis test ~1000 times**
- **(2) Pick the right threshold value such that (if $\alpha = 0.05$) 5% of tests are significant (by chance)**
- **(3) Compare original p-values to new threshold**

Permutation example

You want to test the relationship between genetic variation (across the genome) and worm length. You genotype 10,000 variants and perform a `t.test` for worm length with REF vs. ALT at each variant site.

Question:

Is it statistically significant that worms with the “A” genotype are longer than worms with the “C” genotype?

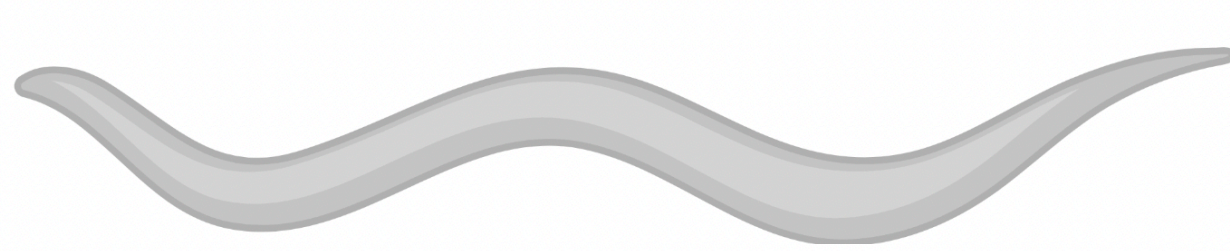
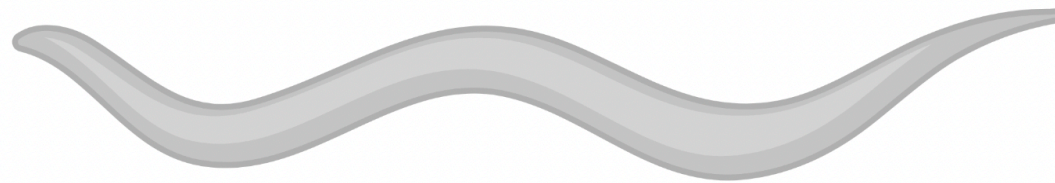
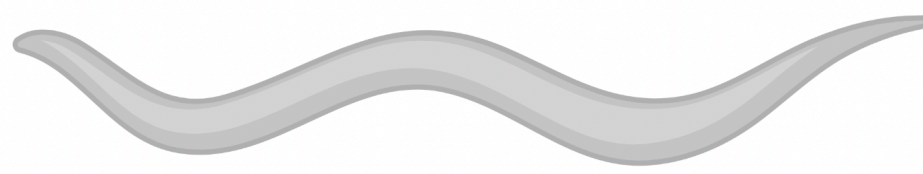
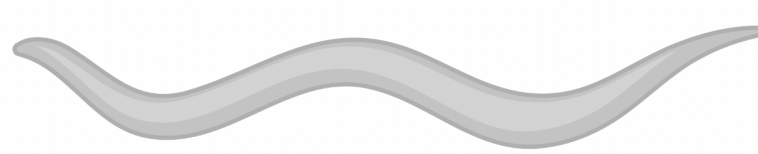
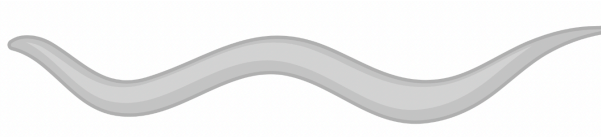

	geno	pheno
	A	450
	A	439
	A	412
	C	400
	C	378
	C	356

Permutation example

You want to test the relationship between genetic variation (across the genome) and worm length. You genotype 10,000 variants and perform a `t.test` for worm length with REF vs. ALT at each variant site.

Question:

Is it statistically significant that worms with the “A” genotype are longer than worms with the “C” genotype?

	geno	pheno
	A	450
	A	439
	A	412
	C	400
	C	378
	C	356



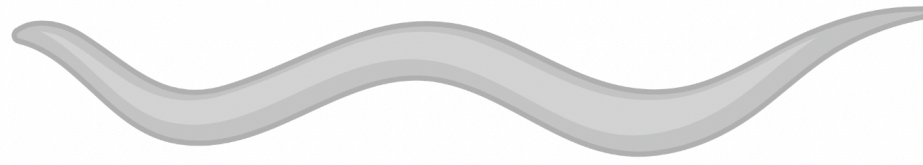
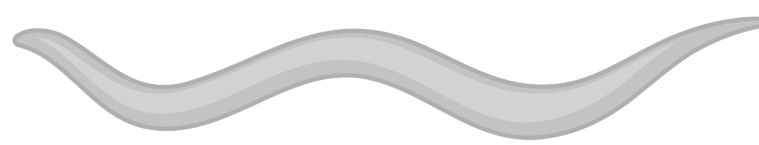


```
> t.test(pheno_A, pheno_C)
```


Permutation example

You want to test the relationship between genetic variation (across the genome) and worm length. You genotype 10,000 variants and perform a `t.test` for worm length with REF vs. ALT at each variant site.

Question:

Is it statistically significant that worms with the “A” genotype are longer than worms with the “C” genotype?

	geno	pheno
	A	450
	C	439
	C	412
	C	400
	A	378
	A	356

```
> t.test(pheno_A, pheno_C)
```

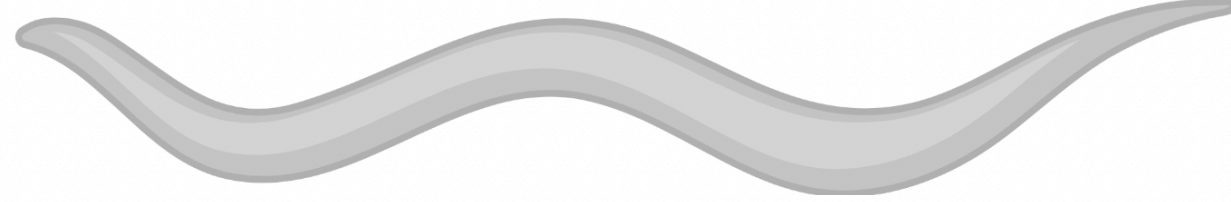
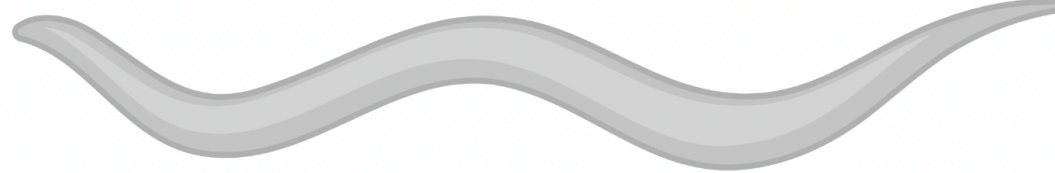
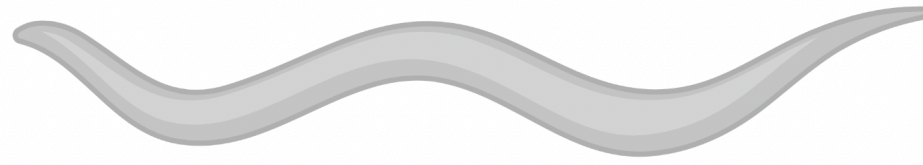
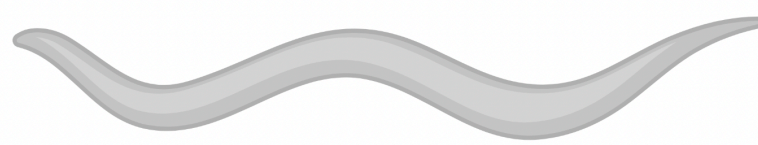


1000x

Permutation example

You want to test the relationship between genetic variation (across the genome) and worm length. You genotype 10,000 variants and perform a `t.test` for worm length with REF vs. ALT at each variant site.

Question:

Is it statistically significant that worms with the “A” genotype are longer than worms with the “C” genotype?

	geno	pheno
	A	450
	C	439
	A	412
	C	400
	C	378
	A	356

```
> t.test(pheno_A, pheno_C)
```

1000x

```
1 #####  
2 # Permutation example #  
3 #####  
4  
5 # create data  
6 set.seed(19)  
7  
8 dist1 <- rnorm(10, -1, 1)  
9 dist2 <- rnorm(10, 1, 1)  
10  
11 # calculate p-value  
12 t.test(dist1, dist2)  
13  
14 # create empty vector for randomization  
15 random_p <- c()  
16  
17 # randomize data
```



Introduction to enrichment analysis

Suppose you measured expression of 10,000 transcripts. After multiple hypothesis correction, you found 200 were significant. Interestingly, many of these 200 significant transcripts looked like they function in the immune system, which would be an exciting discovery!


Should you be excited or skeptical?

Skeptical. Immune genes are relatively abundant in the genome, we need more information. **Do you have more significant immune genes that you'd expect if all gene types were equally abundant?** (i.e. is my list enriched for immune genes?)

Polya urn models and the hypergeometric distribution



Polya urn models and the hypergeometric distribution



**What is the probability
of choosing 2 black
flies and 1 grey fly?**

Is this binomial?

No!!! Why?


(Sampling without
replacement, not
independent,
different p-values)



Polya urn models and the hypergeometric distribution


$\Pr\{2 \text{ black} + 1 \text{ white}\}$

=




7/10

X




6/9

X




3/8

= 0.175




7/10

X




3/9

X




6/8

= 0.175




3/10

X



7/9

X



6/8

= 0.175


= 0.525

7



;

3




Polya urn models and the hypergeometric distribution

$$\Pr\{2 \text{ black} + 1 \text{ white}\} = \frac{(\text{How many ways to get 2 black}) \times (\text{How many ways to get 1 white})}{(\text{How many ways to get 3 flies})}$$

$$= \frac{{}_7C_2 \times {}_3C_1}{{}_{10}C_3}$$


$$= 0.525 \quad \checkmark$$



Polya urn models and the hypergeometric distribution

$$\Pr\{x \text{ black}\} = \frac{{}_m C_x \times {}_n C_{k-x}}{{}_{m+n} C_k}$$

	Black	White	Total
Chosen	x	k-x	k
Not chosen	m-x	n-k+x	m+n-k
Total	m	n	m+n




m = number of black flies
n = number of white flies

k = total number of flies chosen
x = number of black flies chosen

Polya urn models and the hypergeometric distribution

$$\Pr\{x \text{ black}\} = \frac{{}_m C_x \times {}_n C_{k-x}}{{}_{m+n} C_k}$$

	Black	White	Total
Chosen	2	1	3
Not chosen	5	2	7
Total	7	3	10



m = number of black flies
 n = number of white flies


k = total number of flies chosen
 x = number of black flies chosen

Polya urn models and the hypergeometric distribution

$$\Pr\{2 \text{ black} + 1 \text{ white}\} = \frac{{}^7C_2 \times {}^3C_1}{{}^{10}C_3}$$

```
> dhyper(x, m, n, k)
```

	Black	White	Total
Chosen	2	1	3
Not chosen	5	2	7
Total	7	3	10



m = number of black flies
 n = number of white flies


k = total number of flies chosen
 x = number of black flies chosen

Polya urn models and the hypergeometric distribution

$$\Pr\{2 \text{ black} + 1 \text{ white}\} = \frac{{}^7C_2 \times {}^3C_1}{{}^{10}C_3}$$

```
> dhyper(x=2, m=7, n=3, k=3)
```

	Black	White	Total
Chosen	2	1	3
Not chosen	5	2	7
Total	7	3	10




m = number of black flies
 n = number of white flies

k = total number of flies chosen
 x = number of black flies chosen

Polya urn models and the **hypergeometric distribution**

What is the probability of choosing at least one white fly (out of 3)?

$$\begin{aligned}\Pr\{\text{at least 1 white}\} &= \Pr\{1 \text{ white}\} + \Pr\{2 \text{ whites}\} + \Pr\{3 \text{ whites}\} \\ &= 1 - \Pr\{0 \text{ whites}\}\end{aligned}$$



m = number of **white** flies
 n = number of **black** flies

k = total number of flies chosen
 x = number of **white** flies chosen

Polya urn models and the hypergeometric distribution


What is the probability of choosing at least one white fly (out of 3)?

$$\Pr\{\text{at least 1 white}\} = \Pr\{1 \text{ white}\} + \Pr\{2 \text{ whites}\} + \Pr\{3 \text{ whites}\}$$

	Black	White	Total
Chosen	3	0	3
Not chosen	4	3	7
Total	7	3	10

$$= 1 - \Pr\{0 \text{ whites}\}$$

$$= 1 - \frac{{m \mathbf{C}_x} \times {n \mathbf{C}_{k-x}}}{m+n \mathbf{C}_k}$$



m = number of **white** flies
 n = number of **black** flies

k = total number of flies chosen
 x = number of **white** flies chosen

Polya urn models and the hypergeometric distribution

What is the probability of choosing at least one white fly (out of 3)?

$$\Pr\{\text{at least 1 white}\} = \Pr\{1 \text{ white}\} + \Pr\{2 \text{ whites}\} + \Pr\{3 \text{ whites}\}$$


	Black	White	Total
Chosen	3	0	3
Not chosen	4	3	7
Total	7	3	10

$$> 1 - \text{dhyper}(x=0, m=3, n=7, k=3)$$

$$= 1 - \Pr\{0 \text{ whites}\}$$

$$= 1 - \frac{{}^3C_0 \times {}^7C_{3-0}}{{}^{3+7}C_3}$$

$$= 0.708$$



m = number of **white** flies
 n = number of **black** flies

k = total number of flies chosen
 x = number of **white** flies chosen

Polya urn models and the hypergeometric distribution

What is the probability of choosing at least one white fly (out of 3)?

$$\Pr\{\text{at least 1 white}\} = \Pr\{1 \text{ white}\} + \Pr\{2 \text{ whites}\} + \Pr\{3 \text{ whites}\}$$


	Black	White	Total
Chosen	3	0	3
Not chosen	4	3	7
Total	7	3	10

$$> 1 - \text{phyper}(q=0, m=3, n=7, k=3)$$

$$= 1 - \Pr\{0 \text{ whites}\}$$

$$= 1 - \frac{{}^3C_0 \times {}^7C_{3-0}}{{}^{3+7}C_3}$$

$$= 0.708$$



m = number of **white** flies
 n = number of **black** flies

k = total number of flies chosen
 x = number of **white** flies chosen

Polya urn models and the hypergeometric distribution

What is the probability of choosing at least one white fly (out of 3)?

$$\Pr\{\text{at least 1 white}\} = \Pr\{1 \text{ white}\} + \Pr\{2 \text{ whites}\} + \Pr\{3 \text{ whites}\}$$


	Black	White	Total
Chosen	3	0	3
Not chosen	4	3	7
Total	7	3	10

$$= 1 - \Pr\{0 \text{ whites}\}$$

$$= 1 - \frac{{}^3C_0 \times {}^7C_{3-0}}{{}^{3+7}C_3}$$

```
> phyper(q=0, m=3, n=7, k=3, lower.tail = F)
```

= 0.708



m = number of **white** flies
n = number of **black** flies

k = total number of flies chosen
x = number of **white** flies chosen

Introduction to enrichment analysis

Suppose you measured expression of 10,000 transcripts. After multiple hypothesis correction, you found 200 were significant. Interestingly, **80** of these 200 significant transcripts looked like they function in the immune system, which would be an exciting discovery!

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + P(200)$$

	Significant	Not sig.	Total
Immune	80		3,000
Not immune			
Total	200	9,800	10,000

Introduction to enrichment analysis

Suppose you measured expression of 10,000 transcripts. After multiple hypothesis correction, you found 200 were significant. Interestingly, **80** of these 200 significant transcripts looked like they function in the immune system, which would be an exciting discovery!

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + \text{P}(200)$$

	Significant	Not sig.	Total
Immune	80	2,920	3,000
Not immune	120	6,880	7,000
Total	200	9,800	10,000

We could do it by hand...

... but it would take a LONG time

Introduction to enrichment analysis

Suppose you measured expression of 10,000 transcripts. After multiple hypothesis correction, you found 200 were significant. Interestingly, **80** of these 200 significant transcripts looked like they function in the immune system, which would be an exciting discovery!

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + \text{P}(200)$$

	Significant	Not sig.	Total
Immune	80	2,920	3,000
Not immune	120	6,880	7,000
Total	200	9,800	10,000

We could do it by hand...

... but it would take a LONG time

Strictly **GREATER THAN** (not equal to)

```
> phyper(q=79, m=200, n=9800, k=3000, lower.tail = F)
```

Introduction to enrichment analysis

Suppose you measured expression of 10,000 transcripts. After multiple hypothesis correction, you found 200 were significant. Interestingly, **80** of these 200 significant transcripts looked like they function in the immune system, which would be an exciting discovery!

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + P(200)$$

	Significant	Not sig.	Total
Immune	80	2,920	3,000
Not immune	120	6,880	7,000
Total	200	9,800	10,000

[1] 0.001479778

Strictly **GREATER THAN** (not equal to)

```
> phyper(q=79, m=200, n=9800, k=3000, lower.tail = F)
```


Enrichment and the **Fisher's Exact Test**

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + \text{P}(200)$$

	Significant	Not sig.	Total
Immune	80	2,920	3,000
Not immune	120	6,880	7,000
Total	200	9,800	10,000

$$H_0 : P(\text{Sig. immune}) = P(\text{not sig immune})$$

$$H_A : P(\text{Sig. immune}) > P(\text{not sig immune})$$

```
# make data frame
```

```
> genes <- data.frame(sig = c(80, 120), not_sig = c(2920, 6880))
```

```
# run fishers exact test - greater (for enrichment)
```

```
> fisher.test(genes, alternative = "greater")
```


Enrichment and the **Fisher's Exact Test**

$$\text{Pr}\{80+ \text{ sig. immune genes}\} = \text{Pr}(80) + \text{Pr}(81) + \text{Pr}(82) + \dots + \text{P}(200)$$

	Significant	Not sig.
Immune	80	2,920
Not immune	120	6,880
Total	200	9,800

Fisher's Exact Test for Count Data

[1] 0.001479778

data: genes
p-value = 0.00148

alternative hypothesis: true odds ratio is greater than 1
95 percent confidence interval:

1.220544 Inf

sample estimates:

odds ratio
1.5707

```
# make data frame
```

```
> genes <- data.frame(sig
```

```
# run fishers exact test - greater (for enrichment)
```

```
> fisher.test(genes, alternative = "greater")
```

Enrichment and the **Fisher's Exact Test**

Q: Is the Fisher's exact test a parametric or non-parametric test?

A: It is a non-parametric test! (Because there are no assumptions about the underlying population distribution)

Q: Why is it called the Fisher's EXACT test?


A: It calculates the EXACT p -value (the probability that we see OUR data (or more extreme) out of all the possible combinations). Unlike a t -test, the p -value is NOT estimated from a distribution.

Q: Can we have a non-directional Fisher's exact test?

A: Yes! But cumbersome to calculate by hand... almost always want to use R's `fisher.test()`

You are studying the prevalence of a disease in wild field mice. You catch 30 mice (14 southern grasshopper mice and 16 four-striped grass mice) and check for disease. You find that 5 southern grasshopper and 1 four-striped mice have the disease.

Is there a disease enrichment in the southern grasshopper mice?



Southern grasshopper mouse



Four-striped grass mouse

You are studying the prevalence of a disease in wild field mice. You catch 30 mice (14 southern grasshopper mice and 16 four-striped grass mice) and check for disease. You find that 5 southern grasshopper and 1 four-striped mice have the disease.

Is there a disease enrichment in the southern grasshopper mice?

H_0 : P(southern disease) = P(striped disease) H_A : P(southern disease) > P(striped disease)

	Southern grasshopper	Striped grass	Total
Disease	5	1	6
No disease	9	15	24
Total	14	16	30

```
mice <- data.frame(  
  southern = c(5, 9),  
  striped = c(1, 15)  
)
```

```
fisher.test(mice,  
  alternative = "greater")
```

p-value = 0.059

You are studying the prevalence of a disease in wild field mice. You catch 30 mice (14 southern grasshopper mice and 16 four-striped grass mice) and check for disease. You find that 5 southern grasshopper and 1 four-striped mice have the disease.

Is there a disease enrichment in the southern grasshopper mice?

$$P(5+ \text{ southern disease}) = P(5) + P(6)$$

	Southern grasshopper	Striped grass	Total
Disease	5	1	6
No disease	9	15	24
Total	14	16	30

$$> 1 - \text{phyper}(4, 14, 16, 6)$$

$$P(5) = \frac{{}^{14}C_5 \times {}^{16}C_1}{{}^{30}C_6}$$

$$> \text{dhyper}(5, 14, 16, 6)$$

$$P(6) = \frac{{}^{14}C_6 \times {}^{16}C_0}{{}^{30}C_6}$$

$$> \text{dhyper}(6, 14, 16, 6)$$

Announcements

- Thank you for your mid-course survey responses! I will do my best to address issues brought up in the remainder of the course
- No office hours today — will resume Friday @ 1:30 pm
- Practicum #2 on Thursday! **Bring your laptops**
 - Practicum #2 will be an assignment, anything not completed will be homework due **Tuesday, November 9**

Midterm

- If you would like to discuss your exam (or anything else) with me, I will be holding special office hours next week: **Thurs 11/11. 10am - 12pm (Location: TBD)**
 - *If that time doesn't work for you, reach out and we can schedule something else*
- Re-grades will be assessed case-by-case:
 - Must be requested by **Thursday, Nov. 11 3:30 pm**
 - To request: turn in, to Katie, your exam plus thorough written explanation of where/why you feel you deserve more points
 - *Note: all re-grades are subjected to a re-grade of the **entire** exam*

Mean: 88%

Median: 92%

Max: 98%